



Array Computing with Eigen

Jonas Lindemann

jonas.lindemann@naiss.se

Jonas.Lindemann@lunarc.lu.se



Practical information

Schedule

Time	Description	Type
9.00 - 9.05	Introduciton	Lecture
9.05 - 9.15	Setting up development environment	Lab/Interactive
9.15 - 10.00	Working with matrices and vectors	Lecture
10.00 - 10.20	Coffee break	
10.20 - 11.00	Working with exercises	Lab
11.00 - 12.00	Advanced matrix operations	Lecture
12.00 - 13.00	Lunch	
13.00 - 13.40	Working with exercises	Lab
13.40 - 14.00	Best practice and integration	Lecture
14.00 - 14.30	Working with exercises	Lab
14.30 - 14.45	Coffee break	Lecture
14.45 - 15.15	Using Eigen in Parallel applications	Lecture

Course page

- Main page
 - <https://array-computing-with-eigen.readthedocs.io/en/latest/index.html>
- Exercises
 - <https://array-computing-with-eigen.readthedocs.io/en/latest/exercises.html>



Array Computing with Eigen

Contents:

- Setup
- Schedule
- Links
- Exercises

Array Computing with Eigen

[Show Source](#)

This is a course on array computing with the C++ library Eigen. The course is intended for students with a basic knowledge of C++ programming and linear algebra.

Jonas Lindemann, 2025



Introduction

Why Not Build Your Own Matrix Library

- Arrays are crucial for scientific computing applications
- In C++, it might be tempting to create your own matrix library
- However, this is often not recommended because:
 - Existing libraries are well-tested
 - Existing libraries are highly optimized
 - Creating your own introduces unnecessary complexity



What is Eigen

- One of the most popular libraries for linear algebra in C++
- A header-only library (no linking required)
- Fast and easy to use
- Supports optimized packages like BLAS and LAPACK
- Can significantly speed up computational tasks
 - Some operations have support for OpenMP
 - Using optimized libraries will provide additional speed.



Setting up Eigen in Your Project

- Eigen is header-only, so no linking needed
 - Can still need linking for optimized BLAS / LAPACK libraries
- Simply include the relevant header:

```
#include <Eigen/Dense>
```

- The Dense module includes all basic matrix operations
- Most commonly used module for linear algebra tasks
- Other specialized modules are available for specific needs

Getting Started with Eigen

- Easy to integrate into existing C++ projects
- No external dependencies required
- Compatible with modern C++ standards
- Excellent documentation available
- Active community support

Setting Up Your Eigen Development Environment

Environment Setup

- COSMOS at LUNARC (HPC environment)
- Windows
- Linux (Ubuntu)
- macOS

Using COSMOS at LUNARC

- Remote desktop environment: [LUNARC Documentation](#)
- SSH access: [SSH Login Documentation](#)

LUNARC: Loading the Environment

Load required modules:

```
module load foss/2024a
module load Eigen
module load Cmake
```

Verify installation:

```
g++ --version      # Should show GCC 13.3.0
cmake --version     # Should show CMake 3.29.3
```


LUNARC: Testing Eigen

Create a test file `ex0.cpp`:

```
#include <iostream>
#include <Eigen/Dense>

int main()
{
    Eigen::Matrix3d m = Eigen::Matrix3d::Random();
    std::cout << "Here is the matrix m:" << std::endl;
    std::cout << m << std::endl;
    return 0;
}
```

LUNARC: Compiling and Running

Compile:

```
g++ ex0.cpp -o ex0
```

Run:

```
./ex0
```

Expected output:

Here is the matrix m:

```
0.680375    0.59688 -0.329554
```

```
-0.211234   0.823295  0.536459
```

```
0.566198  -0.604897 -0.444451
```

Ubuntu Linux Setup

Install required packages:

```
sudo apt-get update
sudo apt-get install g++
sudo apt-get install cmake
sudo apt-get install libeigen3-dev
```

Verify installation:

```
g++ --version
cmake --version
```

Ubuntu Linux: Testing Eigen

Create the same test file `ex0.cpp` as before.

Compile with include path:

```
g++ ex0.cpp -I/usr/include/eigen3 -o ex0
```

Run:

```
./ex0
```

The output should match the previous example.

macOS Setup

Install required packages using Homebrew:

```
brew install gcc  
brew install cmake  
brew install eigen
```

macOS: Testing Eigen

Create the same test file `ex0.cpp` as before.

Compile with include path:

```
g++ -std=c++11 -I/usr/local/include/eigen3 ex0.cpp -o ex0
```

Run:

```
./ex0
```

The output should match the previous examples.

Troubleshooting Tips

- Check compiler version compatibility
- Verify include paths are correct
- Ensure Eigen headers are properly installed
- For library issues, confirm environment variables are set correctly



Demo

A close-up, slightly angled photograph of a traditional wooden abacus. The abacus has a light-colored wooden frame and several horizontal rods. Each rod is populated with dark, oval-shaped beads. A single pencil, with a sharpened lead tip, lies diagonally across the middle of the abacus, resting on one of the rods. The background is a dark, solid color, providing a high contrast for the light-colored abacus.

Working with Matrices and Vectors

Matrix and Vector Types

- All Eigen classes are template classes (work with different data types)
- Most common data types: **float**, **double**, and **int**
- All classes defined in the **Eigen** namespace
- Best practice: use **Eigen::** prefix instead of
 - **using namespace Eigen**
- In many of my example I will use without the prefix to make the code more readable.

Matrix Declaration

Generic form:

```
// 3x3 matrix of doubles
```

```
Eigen::Matrix<double, 3, 3> A;
```

Using convenient typedefs:

```
Eigen::Matrix3d B; // Same as above
```

Vector Declaration

Generic form:

```
// 3x1 vector of doubles
```

```
Eigen::Vector<double, 3> v;
```

Using convenient typedefs:

```
Eigen::Vector3d w; // Same as above
```

Note: Vectors are matrices with one dimension fixed to 1

Initialization

- Newly declared matrices contain **random values**
- Methods to initialize:
 - `A.setZero()` - Initialize to zeros
 - `A.setOnes()` - Initialize to ones
 - `<< operator` - Set specific values

```
Eigen::Matrix3d D;
```

```
D << 1, 2, 3,  
      4, 5, 6,  
      7, 8, 9;
```

Vectors - Special Properties

- Initialize with constructor:

```
Eigen::Vector3d v(1, 2, 3);
```

- Initialize with the << operator:

```
Eigen::Vector3d w;  
w << 1, 2, 3;
```

- Special initialization methods:

```
v.setLinSpaced(3, 1, 2); // Creates: [1, 1.5, 2]  
w.setRandom();           // Random values
```

Column and Row Access

// Column insertion

```
D.col(0) << 1, 2, 3;  
D.col(1) << 4, 5, 6;  
D.col(2) << 7, 8, 9;
```

// Row insertion

```
D.row(0) << 1, 2, 3;  
D.row(1) << 4, 5, 6;  
D.row(2) << 7, 8, 9;
```

Common Typedefs

Fixed-size matrices:

```
Matrix2d, Matrix3d, Matrix4d  // double  
Matrix2f, Matrix3f, Matrix4f  // float  
Matrix2i, Matrix3i, Matrix4i  // int
```

Dynamic-size matrices:

```
MatrixXd  // double  
MatrixXf  // float  
MatrixXi  // int
```


Common Vector Typedefs

Fixed-size vectors:

```
Vector2d, Vector3d, Vector4d // double  
Vector2f, Vector3f, Vector4f // float  
Vector2i, Vector3i, Vector4i // int
```

Dynamic-size vectors:

```
VectorXd // double  
VectorXf // float  
VectorXi // int
```



Demo

Fixed vs Dynamic Size

Fixed size (known at compile time):

```
// 3x3 matrix, size fixed at compile time
```

```
Eigen::Matrix3d A;
```

Dynamic size (determined at runtime):

```
Eigen::Matrix<double, Eigen::Dynamic, Eigen::Dynamic> F;  
F.resize(3, 3); // Set dimensions at runtime
```

```
// Shorthand
```

```
Eigen::MatrixXd G(3, 3);
```

Resizing Matrices

- `resize(rows, cols)` changes matrix dimensions
- If the total number of elements stays the same, data is preserved
- If the total elements change, data is lost and must be reinitialized

```
A_dyn.resize(1, 9); // Preserves data if A_dyn was 3x3  
A_dyn.resize(6, 6); // Changes total elements, data is lost
```

Row Vectors

// 1x3 vector

```
Eigen::RowVector3d r(1.0, 2.0, 3.0);
```

// Output: 1 2 3

// 3x1 vector

```
Eigen::Vector3d s(1.0, 2.0, 3.0);
```

// Output:

// 1

// 2

// 3

Matrix Operations

- Addition: $A + B$
- Scalar multiplication: $A * 3.0$
- Matrix multiplication: $A * B$
- Add scalar to all elements: $E + \text{Matrix3d}::\text{Constant}(1.0)$
- Element-wise operations: $E.\text{array}() + 3.0$

Matrix Transformations

- Transpose: `A.transpose()`
- Inverse: `A.inverse()`
- Component-wise operations:

```
Vector3d x(1, 4, 9);
```

```
auto y = x.cwiseSqrt(); // [1, 2, 3]
```

```
// Or
```

```
auto w = z.array().sqrt();
```

Vector Operations

- Dot product: `s.dot(t)`
- Cross product: `s.cross(t)`
- Component-wise operations: `x.cwiseSqrt()`

Reduction Operations

- `K.sum()` - Sum of all elements
- `K.prod()` - Product of all elements
- `K.mean()` - Mean of all elements
- `K.norm()` - Euclidean norm
- `K.maxCoeff()` - Maximum value
- `K.minCoeff()` - Minimum value
- `K.trace()` - Trace of the matrix
- `K.diagonal()` - Diagonal elements
- `K.determinant()` - Determinant



Demo

A photograph of a coffee break. In the foreground, a black cup filled with a frothy coffee sits on a matching black saucer. To the left, a white plate holds a golden-brown croissant. The entire scene is set on a light-colored wooden table. The text "Coffee break" is overlaid in white, centered horizontally. Below it, the time "10:00 – 10:20" is also overlaid in white.

Coffee break

10:00 – 10:20

The background of the slide is a top-down view of numerous LEGO bricks of various colors (red, blue, yellow, green, orange) and shapes (1x1, 1x2, 1x3, 2x2, 2x4, etc.) scattered across a solid yellow surface. The bricks are arranged in a seemingly random pattern, with some showing their studs and others showing their smooth sides. The lighting is even, highlighting the texture of the plastic.

Exercise

10:20 – 11:00

Advanced Matrix Operations in Eigen

Reshaping Matrices

- `.reshaped(rows, cols)` method allows changing matrix dimensions
 - Returns a view into the original matrix (not a copy)
 - Changes to the reshaped matrix affect the original
- `Matrix3d A;`

```
A << 1, 2, 3,  
      4, 5, 6,  
      7, 8, 9;
```

```
auto B = A.reshaped(1, 9); // Result: 1 2 3 4 5 6 7 8 9
```

Reshaping Considerations

- Data is stored in column-major order in Eigen
- Self-assignment requires `.eval()` to force evaluation:
`C = C.reshapeed(1, 9).eval();`
- Can combine with `.transpose()`:
`MatrixXd D = C.reshapeed(1, 9).transpose();`



Demo

Slicing and Indexing

Row and column access:

```
// Set values in row 3
```

```
A.row(3) << 1, 2, 3, 4, 5, 6, 7, 8, 9, 10;
```

```
// Set values in column 3
```

```
A.col(3) << 1, 2, 3, 4, 5, 6, 7, 8, 9, 10;
```

```
// Set all elements in column 1 to 1
```

```
A.col(1).setOnes();
```


Range Selection with `seq()`

Select ranges of rows and columns:

```
// Select block from rows 3-5, columns 3-5
```

```
B(seq(3, 5), seq(3, 5)).setConstant(1);
```

```
// Select every other row/column from 0-9
```

```
B(seq(0, 9, 2), seq(0, 9, 2)).setConstant(2);
```

Special Selectors

// Set all elements in the last column to 3

```
B(all, last).setConstant(3);
```

*// Set all elements in the second-to-last
// column to 4*

```
B(all, last - 1).setConstant(4);
```


Index-Based Selection

Use `std::vector` of indices to select specific rows and columns:

```
vector<int> idx = { 1, 3, 4, 6, 7, 9 };
```

```
// Select submatrix using index vector
```

```
auto D = C(idx, idx);
```



Demo

Linear System Solving - Small Matrices

For small matrices (up to 4x4), using inverse is acceptable:

```
Matrix3d A;  
A.setRandom();  
Vector3d b;  
b.setRandom();
```

```
// Solve  $Ax = b$   
Vector3d x = A.inverse() * b;
```

Linear System Solving - Larger Matrices

For larger matrices, use decomposition methods:

```
MatrixXd A(10, 10);  
A.setRandom();  
VectorXd b(10);  
b.setRandom();
```

```
// Solve using QR decomposition
```

```
VectorXd x = A.colPivHouseholderQr().solve(b);
```

```
// Check error
```

```
double error = (A * x - b).norm();
```



Demo

Matrix Decompositions

Different decompositions for different matrix types:

- `colPivHouseholderQr()` - General matrices (robust)
- `fullPivLu()` - General matrices (most stable, slower)
- `ldlt()` - Symmetric matrices
- `householderQr()` - General matrices (fastest, less accurate)

Reusing Decompositions

Create decomposition objects for reuse with multiple right-hand sides:

```
// Create decomposition once  
FullPivLU<MatrixXd> lu(A);
```

```
// Solve for multiple right-hand sides  
VectorXd x1 = lu.solve(b1);  
VectorXd x2 = lu.solve(b2);
```

```
// Solve systems with multiple right-hand sides at once  
MatrixXd X = lu.solve(B); // B has multiple columns
```




Demo



Lunch

12:00 – 13:00

An illustration of a person with short, wavy brown hair and glasses, wearing a light blue long-sleeved shirt and dark trousers, standing in a workspace. They are looking at a large computer monitor that displays various data visualizations, including bar charts, line graphs, and a map. The person's right hand is near the monitor. The background is a wall covered with numerous colorful sticky notes and papers. To the right, there is a desk with a lamp, a stack of papers, and other office supplies. The overall style is a soft, painterly illustration with a warm, slightly muted color palette.

Eigen Best Practices and Integration

Returning Matrices from Functions

- Prefer returning Eigen matrices by value
- C++ return value optimization (RVO) prevents unnecessary copying
- Simple and effective approach:

```
MatrixXd foo()
{
    MatrixXd A(10, 10);
    A.setRandom();
    return A; // Efficient due to RVO
}

int main()
{
    MatrixXd B = foo(); // No unnecessary copying
}
```

Passing Matrices to Functions

- Pass matrices as const references to avoid copying
- Use the `const` keyword to indicate the matrix won't be modified

```
void bar(const MatrixXd& A)  // Pass by reference
{
    cout << A << endl;
}

int main()
{
    MatrixXd B(10, 10);
    B.setRandom();
    bar(B);  // No copying occurs
}
```

Implementing Functions with Eigen

General Rule: - Pass Eigen matrices and vectors by reference -
Return Eigen matrices and vectors by value

```
// Example: Function that creates and returns a matrix  
MatrixXd hooke(TAnalysisType ptype, double E, double v)  
{  
    MatrixXd D;  
    // ... matrix construction ...  
    return D; // Return by value  
}
```

Example: Function Implementation

```
Matrix4d bar2e(const Vector2d& ex, const Vector2d& ey, const
Vector2d& ep)
{
    // Parameters passed by const reference
    double E = ep(0);
    double A = ep(1);
    double L = sqrt(pow(ex(1)-ex(0),2)+pow(ey(1)-ey(0),2));

    // ... calculations ...

    Matrix4d Ke = G.transpose()*Ke_loc*G;
    return Ke; // Return by value
}
```

Accessing Raw Data

- Use `.data()` method to get pointer to raw data
- Useful for interfacing with C-style APIs

```
MatrixXd A(10, 10);  
A.setRandom();  
  
// Get pointer to underlying data  
double* data = A.data();  
  
// Access elements directly  
for (int i = 0; i < A.size(); i++) {  
    cout << data[i] << " ";  
}
```


Creating 2D Array Views

- Eigen stores data as 1D array internally
- Create array of pointers for 2D access:

```
double* data = A.data();  
double** data2D = new double*[A.rows()];  
  
// Setup row pointers  
for (int i = 0; i < A.rows(); i++)  
    data2D[i] = data + i * A.cols();  
  
// Clean up (only deletes pointer array, not actual data)  
delete[] data2D;
```

Integrating with Other Libraries

- When working with libraries requiring 2D C-style arrays:

```
void foo(double** data, int rows, int cols) {  
    // Function expecting 2D C-style array  
}
```

```
// Setup for integration  
double** data2D = new double*[A.rows()];  
for (int i = 0; i < A.rows(); i++)  
    data2D[i] = A.row(i).data();
```

```
// Call external function  
foo(data2D, A.rows(), A.cols());
```

```
// Clean up  
delete[] data2D;
```

Dealing with Const Issues

- Use `const_cast` when necessary for C-style API integration:

```
// If compiler warns about const correctness  
data2D[i] = const_cast<double*>(A.row(i).data());
```

- Only do this when sure the data won't be modified
- Avoid if possible - const correctness is a feature, not a bug

Summary of Best Practices

1. Return matrices by value
2. Pass matrices by const reference
3. Use `.data()` for raw data access
4. Create proper interfaces for C-style libraries
5. Maintain const correctness when possible
6. Be aware of memory ownership when interfacing



Demo

The background of the slide is a top-down view of numerous LEGO bricks of various colors (red, blue, yellow, green, orange) and shapes (1x1, 1x2, 1x3, 2x2, 2x4, etc.) scattered across a solid yellow surface. The bricks are arranged in a seemingly random pattern, with some showing their studs and others showing their hollows or unique shapes like Technic connectors.

Exercise

10:20 – 11:00

A photograph of a coffee break. In the foreground, a black cup filled with a frothy coffee sits on a matching black saucer. To the left, a white plate holds a golden-brown croissant. The entire scene is set on a light-colored wooden table. The text "Coffee break" is overlaid in white, centered horizontally. Below it, the time "10:00 – 10:20" is also overlaid in white.

Coffee break

10:00 – 10:20

A photograph of three model trains on a track at night. The train on the left is red and white, the middle one is yellow and white, and the one on the right is blue. They are all facing forward with their headlights on. The background is dark with some blurred lights.

Using Eigen in Parallel Applications

Overview

- Eigen provides efficient array and vector data types
- Takes advantage of underlying linear algebra libraries if available
- Raw data accessible through `.data()` method
- Many operations support OpenMP parallelization

Eigen and OpenMP Support

Eigen operations that use OpenMP for parallelization:

- general dense matrix - matrix products
- PartialPivLU
- row-major-sparse * dense vector/matrix products
- ConjugateGradient with Lower|Upper as the UpLo template parameter.
- BiCGSTAB with a row-major sparse matrix format.
- LeastSquaresConjugateGradient

Using Eigen with OpenMP

Simple example: Matrix multiplication using Eigen's built-in OpenMP support

```
// Setting number of threads  
omp_set_num_threads(n_threads);  
  
// This operation will automatically use OpenMP  
MatrixXd C = A * B;
```


OpenMP Performance Scaling

Performance scaling for matrix
multiplication (20000×20000):

Threads	Time (seconds)
1	870.111
2	427.386
4	217.594
8	109.546
12	74.313
24	40.574
48	22.657

Custom Parallelization (Part 1)

Using Eigen for storage with custom OpenMP implementation:

```
// Use RowMajor for more efficient row operations
using Matrix = Eigen::Matrix<double, Eigen::Dynamic,
                             Eigen::Dynamic, Eigen::RowMajor>;
using Vector = Eigen::VectorXd;

Vector customMatVecMult(const Matrix& A, const Vector& x) {
    const int rows = A.rows();
    const int cols = A.cols();
    Vector result(rows);
```

Custom Parallelization (Part 2)

```
// Get raw pointers to the data
const double* A_data = A.data();
const double* x_data = x.data();
double* result_data = result.data();

#pragma omp parallel for schedule(static)
for (int i = 0; i < rows; i++) {
    double sum = 0.0;
    const double* row = A_data + i * cols;

    for (int j = 0; j < cols; j++) {
        sum += row[j] * x_data[j];
    }

    result_data[i] = sum;
}

return result;
}
```


Performance Comparison

Custom OpenMP implementation vs. Eigen's built-in
(40000×40000):

Matrix size: 40000x40000

OpenMP implementation time: 352ms

Eigen implementation time: 389ms

Relative error: 7.35186e-15

Custom implementation slightly outperforms Eigen's built-in operation.

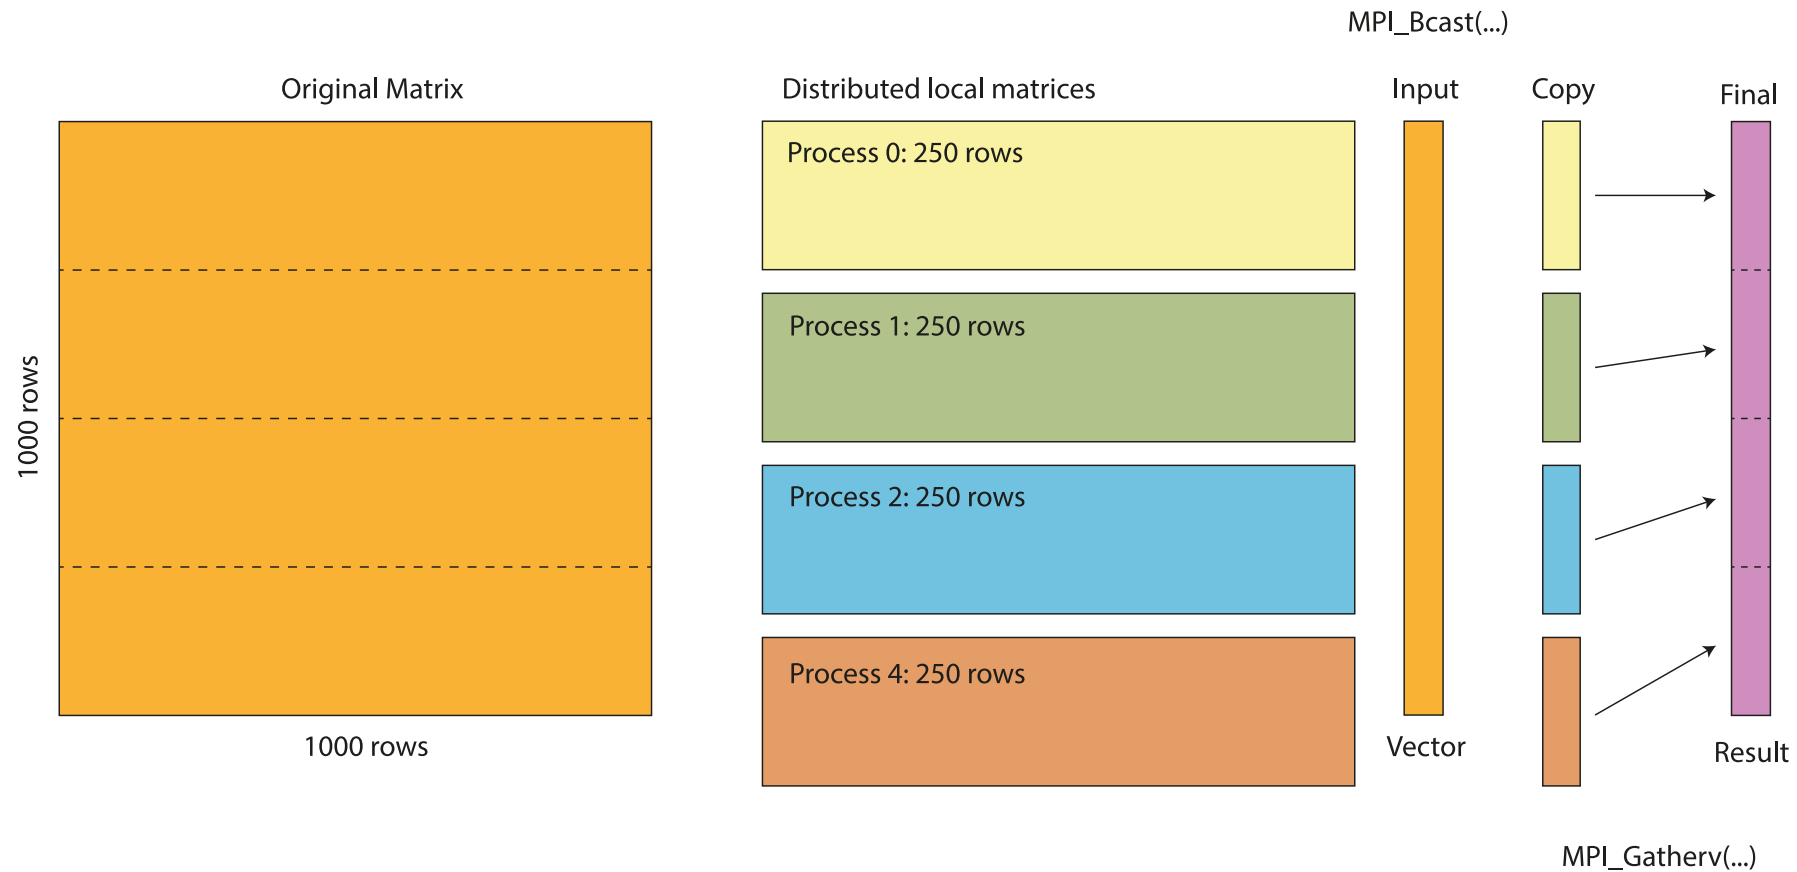
Using Eigen with MPI

- Eigen is NOT a distributed Matrix library
- MPI for distributed memory parallelism
- Eigen arrays stored in 1D memory layout
- Developer handles data distribution
- Example: matrix-vector multiplication across multiple processes

MPI Distribution Strategy

- Matrix divided by rows
- Each process handles a portion of rows
- Vector broadcasted to all processes
- Each process computes partial result

MPI Distribution Strategy



MPIMatrix Class (Part 1)

```
class MPIMatrix {  
private:  
    int m_rank;  
    int m_size;  
    int m_rows;  
    int m_cols;  
    MatrixXd m_localMatrix;  
    VectorXd m_localResult;  
  
public:  
    MPIMatrix(int r, int c)  
        : m_rank(0), m_size(1), m_rows(r), m_cols(c)  
    {  
        MPI_Comm_rank(MPI_COMM_WORLD, &m_rank);  
        MPI_Comm_size(MPI_COMM_WORLD, &m_size);  
    }  
};
```

MPIMatrix Class (Part 2)

```
// Calculate local matrix size
int localRows = m_rows / m_size;
if (m_rank < m_rows % m_size) {
    localRows++;
}

m_localMatrix.resize(localRows, m_cols);
m_localResult.resize(localRows);
}

void randomize() {
    srand(std::chrono::system_clock::now().time_since_epoch().count());
    m_localMatrix.setRandom();
}
```

MPIMatrix Class (Part 3)

```
void multiply(const VectorXd& vec) {  
    // Local multiplication  
    m_localResult = m_localMatrix * vec;  
}  
  
void printResult() const {  
    // Gather results  
    std::vector<int> recvCounts(m_size);  
    std::vector<int> displs(m_size);  
  
    // Calculate receive counts and displacements  
    int baseCount = m_rows / m_size;  
    int remainder = m_rows % m_size;
```


MPIMatrix Class (Part 4)

```
for (int i = 0; i < m_size; ++i) {
    recvCounts[i] = baseCount + (i < remainder ? 1 : 0);
    displs[i] = (i > 0) ? displs[i-1] + recvCounts[i-1] : 0;
}

// Allocate space for complete result
VectorXd globalResult;
if (m_rank == 0)
    globalResult.resize(m_rows);

// Gather all local results to rank 0
MPI_Gatherv(m_localResult.data(), m_localResult.size(), MPI_DOUBLE,
            globalResult.data(), recvCounts.data(), displs.data(),
            MPI_DOUBLE, 0, MPI_COMM_WORLD);
```

MPIMatrix Class (Part 5)

```
// Print result on rank 0
if (m_rank == 0) {
    std::cout << "First few elements of result: \n"
               << globalResult.head(5).transpose() << std::endl;
}
};
```

MPI Main Function (Part 1)

```
int main(int argc, char** argv) {  
    constexpr int MatrixSize = 10000;  
    int rank;  
  
    MPI_Init(&argc, &argv);  
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);  
  
    // Create distributed matrix  
    MPIMatrix distMatrix(MatrixSize, MatrixSize);  
    distMatrix.randomize();  
}
```

MPI Main Function (Part 2)

```
// Create a random x vector
VectorXd x;
if (rank == 0) {
    std::cout << "Generating random vector x...\n";
    x = VectorXd::Random(MatrixSize);
} else {
    x.resize(MatrixSize);
}

// Broadcast x vector to all processes
MPI_Bcast(x.data(), MatrixSize, MPI_DOUBLE, 0, MPI_COMM_WORLD);
```

MPI Main Function (Part 3)

```
// Perform distributed matrix-vector multiplication
auto startTime = std::chrono::high_resolution_clock::now();
distMatrix.multiply(x);
auto endTime = std::chrono::high_resolution_clock::now();

if (rank == 0) {
    auto duration = std::chrono::duration_cast<std::chrono::milliseconds>(endTime-startTime);
    std::cout << "Matrix size: " << MatrixSize << " x " << MatrixSize << std::endl;
    std::cout << "Matrix memory size (MB): "
                << sizeof(double) * MatrixSize * MatrixSize / 1e6 << std::endl;
    std::cout << "Multiplication completed in " << duration.count() << " ms\n";
}

distMatrix.printResult();
```

MPI Considerations

- For large datasets, avoid gathering all data to one process
- Consider:
 - Writing results to separate files on each rank
 - Using parallel I/O libraries like HDF5
 - Using MPI I/O functions

Summary: Eigen in Parallel Applications

- Eigen works well with OpenMP and MPI
- Use built-in Eigen parallel operations when possible
- For custom parallelization:
 - Access raw data with `.data()`
 - Use RowMajor order for row-wise operations
 - Implement your own parallelization with OpenMP/MPI
- Benefits:
 - Simplified memory management
 - Clean, maintainable code
 - Good performance



Questions / Answers